


Jev is the fastest option in this relation-selection run.

431 ms p95

It also has the lowest planning-cost signal. Accuracy is lower than both chat comparators, so this is an operational tradeoff - not a general model ranking.



Exactly what was measured - and what was not.

Each provider saw the same one-case request, 16 relation descriptions, and an opaque case ID. Failed calls remain in the denominator and release evidence.

INPUT	CHOICE	EXECUTION	SCORE
One positive case Sentence plus supplied source and target entities.	16 relations No abstention in this positive-only sample.	160 requests / model Sequential, temperature 0, 55 s timeout, no retry.	Planned-case accuracy Uncovered cases count incorrect; gold IDs stay local.

Failure ledger

Jev	HTTP 503 at offset 98; unknown cost; no retry; continuation at 99.
DeepSeek	prediction_outside_criteria; known cost; receipt retained; no retry.
Luna	160 / 160 valid scored decisions.

Ready to publish as a scoped benchmark.

The task, revision, seed, request controls, cost basis, failure receipts, and checksums are all preserved.

LIMITS

- Relation identification only; not end-to-end graph construction.
- Positive-only public sample; possible training-data exposure.
- One seed; no confidence interval or statistical test.
- Jev cost is illustrative, not an invoice-to-invoice comparison.